

# Chapter 17. Deploying Spark

infrastructure you need in place for you and your team to be able to run

- Spark Applications:
- Cluster deployment choices
- Spark's different cluster managers

## **Deployment considerations and configuring deployments**

Spark has three officially supported cluster managers:

- Standalone mode
- Hadoop YARN
- Apache Mesos

These cluster managers maintain a set of machines onto which you can deploy Spark Applications.

## **On-Premises Cluster Deployments**

- reasonable option, especially for organizations that already manage their own datacenters
- on-premises cluster gives you full control over the hardware used, meaning you can optimize performance for your specific workload
- your cluster is fixed in size, whereas the resource demands of data analytics workloads are often elastic
- if you make your cluster too small, it will be hard to launch the occasional very large analytics query or training job for a new machine learning model,
- whereas if you make it large, you will have resources sitting idle
- Second, for on-premises clusters, you need to select and operate your own storage system, such as a Hadoop file system or scalable key-value store
- This includes setting up georeplication and disaster recovery if required
- If you are going to deploy on-premises, the best way to combat the resource utilization problem is to use a cluster manager that allows you to run many Spark applications and dynamically reassign resources between them, or even allows non-Spark applications on the same cluster
- All of Spark's supported cluster managers allow multiple concurrent applications, but YARN and Mesos have better support for dynamic sharing and also additionally support non-Spark workloads
- Handling resource sharing is likely going to be the biggest difference your users see day to day with Spark on-premise versus in the cloud: in public clouds, it's easy to give each application its own cluster of exactly the required size for just the duration of that job
- For storage, you have several different options, but covering all the trade-offs and operational details in depth would probably require its own book
- The most common storage systems used for Spark are distributed file systems such as Hadoop's HDFS and key-value stores such as Apache Cassandra
- Streaming message bus systems such as Apache Kafka are also often used for ingesting data

## Cloud Deploy

- early big data systems were designed for on-premises deployment
- cloud is now an increasingly common platform for deploying Spark
- The public cloud has several advantages when it comes to big data workloads
- resources can be launched and shut down elastically
- can run that occasional “monster” job that takes hundreds of machines for a few hours without having to pay for them all the time
- Even for normal operation, you can choose a different type of machine and cluster size for each application to optimize its cost performance—
  - launch machines with Graphics Processing Units (GPUs) just for your deep learning jobs
- Second, public clouds include low-cost, georeplicated storage that makes it easier to manage large amounts of data
- All the major cloud providers (Amazon Web Services [AWS], Microsoft Azure, Google Cloud Platform [GCP], and IBM Bluemix) include managed Hadoop clusters for their customers, which provide HDFS for storage as well as Apache Spark
- This is actually not a great way to run Spark in the cloud, however, because by using a fixed-size cluster and file system, you are not going to be able to take advantage of elasticity
- Instead, it is generally a better idea to **use global storage systems that are decoupled from a specific cluster**, such as Amazon S3, Azure Blob Storage, or Google Cloud Storage and spin up machines dynamically for each Spark workload
- With decoupled compute and storage, you will be able to pay for computing resources only when needed, scale them up dynamically, and mix different hardware types
- Basically, keep in mind that running Spark in the cloud need not mean migrating an on-premises installation to virtual machines: you can run Spark natively against cloud storage to take full advantage of the cloud’s elasticity, cost-saving benefit, and management tools without having to manage an on-premise computing stack within your cloud environment

Several companies provide “cloud-native” Spark-based services, and all installations of Apache Spark can of course connect to cloud storage. **Databricks, the company started by the Spark team from UC Berkeley, is one example of a service provider built specifically for Spark in the cloud. Databricks provides a simple way to run Spark workloads without the heavy baggage of a Hadoop installation.** The company provides a number of features for running Spark more efficiently in the cloud, such as auto-scaling, auto-termination of clusters, and optimized connectors to cloud storage, as well as a collaborative environment for working on notebooks and standalone jobs. The company also provides a free Community Edition for learning Spark where you can run notebooks on a small cluster and share them live with others. A fun fact is that this entire book was written using the free Community Edition of Databricks, because we found the integrated Spark notebooks, live collaboration, and cluster management the easiest way to produce and test this content.

If you run Spark in the cloud, much of the content in this chapter might not be relevant because you can often create a separate, short-lived Spark cluster for each job you execute. In that case, the standalone cluster manager is likely the easiest to use. However, you may still want to read this content if you'd like to share a longer-lived cluster among many applications, or to install Spark on virtual machines yourself.

## Standalone Mode

- Spark's standalone cluster manager is a lightweight platform built specifically for Apache Spark workloads
- can run multiple Spark Applications on the same cluster
- provides simple interfaces for doing so but can scale to large Spark workloads
- main disadvantage of the standalone mode is that it's more limited than the other cluster managers—in particular, your cluster can only run Spark
- It's probably the best starting point if you just want to

quickly get Spark running on a cluster, however, and you do not have experience using YARN or Mesos

## Starting a standalone cluster

- Starting a standalone cluster requires provisioning the machines for doing so
- starting them up, ensuring that they can talk to one another over the network, and getting the version of Spark you would like to run on those sets of machines
- two ways to start the cluster:
  - by hand
  - using built-in launch scripts.

## launch a cluster by hand

- The first step is to start the master process on the machine that we want that to run on, using the following command:

```
$SPARK_HOME/sbin/start-master.sh
```

- the cluster manager master process will start up on that machine.
- the master prints out a `spark://HOST:PORT` URI
- You use this when you start each of the worker nodes of the cluster, and you can use it as the master argument to your `SparkSession` on application initialization
- You can also find this URI on the master's web UI, which is <http://master-ip-address:8080> by default.
- With that URI, start the worker nodes by logging in to each machine and running the following script using the URI you just received from the master node
- The master machine must be available on the network of the worker nodes you are using, and the port must be open on the master node, as well:

```
$SPARK_HOME/sbin/start-worker.sh <master-spark-URI>
```

## Launch Scripts

- configure cluster launch scripts that can automate the launch of standalone clusters
- create a file called `conf/workers` in your Spark directory that will contain the hostnames of all the machines on which you intend to start Spark workers, one per line
- If this file does not exist, everything will launch locally
- When you go to actually start the cluster, the master machine will access each of the worker machines via Secure Shell (SSH)
- By default, SSH is run in parallel and requires that you configure password-less (using a private key) access
- If you do not have a password-less setup, you can set the environment variable **SPARK\_SSH\_FOREGROUND** and serially provide a password for each worker
- After you set up this file, you can launch or stop your cluster by using the following shell scripts, based on Hadoop's deploy scripts, and available in `$SPARK_HOME/sbin`:

Starts a master instance on the machine on which the script is executed.

```
$SPARK_HOME/sbin/start-master.sh
```

Starts a worker instance on each machine specified in the `conf/workers` file.

```
$SPARK_HOME/sbin/start-workers.sh
```

Starts a worker instance on the machine on which the script is executed.

```
$SPARK_HOME/sbin/start-worker.sh
```

Starts both a master and a number of workers as described earlier.

```
$SPARK_HOME/sbin/start-all.sh
```

Stops the master that was started via the `bin/start-master.sh` script.

```
$SPARK_HOME/sbin/stop-master.sh
```

Stops all worker instances on the machines specified in the `conf/workers` file.

```
$SPARK_HOME/sbin/stop-workers.sh
```

Stops both the master and the workers as described earlier.

```
$SPARK_HOME/sbin/stop-all.sh
```

## Standalone cluster configurations

- Standalone clusters have a number of configurations that you can use to tune your application

- These control everything from what happens to old files on each worker for terminated applications to the worker's core and memory resources
- These are controlled via environment variables or via application properties
- Due to space limitations, we cannot include the entire configuration set here
- Refer to the relevant table on Standalone Environment Variables in the Spark documentation

### Submitting applications

- you can submit applications to it using the `spark://` URI of the master
- you can do this either on the master node itself or another machine using `spark-submit`

## YARN

- Hadoop YARN is a framework for job scheduling and cluster resource management
- Even though Spark is often (mis)classified as a part of the “Hadoop Ecosystem,” in reality, Spark has little to do with Hadoop
- Spark does natively support the Hadoop YARN cluster manager but it requires nothing from Hadoop itself
- You can run your Spark jobs on Hadoop YARN by specifying the master as YARN in the `spark-submit` command-line arguments
- Just like with standalone mode, there are a number of knobs that you are able to tune according to what you would like the cluster to do
- The number of knobs is naturally larger than that of Spark's standalone mode because Hadoop YARN is a generic scheduler for a large number of different execution frameworks

### Submitting applications

- submitting applications to YARN, the core difference from other deployments is that `--master` will become `yarn` as opposed the master node IP, as it is in standalone mode
- Spark will find the YARN configuration files using the environment variable `HADOOP_CONF_DIR` or `YARN_CONF_DIR`

## NOTE

There are two deployment modes that you can use to launch Spark on YARN:

- **cluster mode** has the spark driver as a process managed by the YARN cluster, and the client can exit after creating the application
- **client mode**, the driver will run in the client process and therefore YARN will be responsible only for granting executor resources to the application, not maintaining the master node
  - in cluster mode, Spark doesn't necessarily run on the same machine on which you're executing
  - libraries and external jars must be distributed manually or through the `--jars` command-line argument.

### Configuring Spark on YARN Applications

- variety of different configurations and their implications for your Spark applications

## Hadoop configurations

- HDFS using Spark, you need to include two Hadoop configuration files on Spark's classpath:
  - `hdfs-site.xml` - provides default behaviors for the HDFS client;
  - `core-site.xml` - sets the default file system name
- The location of these configuration files varies across Hadoop versions, but a common location is inside of `/etc/hadoop/conf`
- To make these files visible to Spark, set **HADOOP\_CONF\_DIR** in `$SPARK_HOME/spark-env.sh` to a location containing the configuration files or as an environment variable when you go to spark submit your application

## Apache Mesos

- Apache Mesos is another clustering system that Spark can run on
- the project was also started by many of the original authors of Spark, including one of the authors of this book
- Apache Mesos abstracts CPU, memory, storage, and other compute resources away from machines (physical or virtual), enabling fault-tolerant and elastic distributed systems to easily be built and run effectively
- Mesos intends to be a datacenter scale-cluster manager that manages not just short-lived applications like Spark, but long-running applications like web applications or other resource interfaces
- Mesos is the heaviest-weight cluster manager, simply because you might choose this cluster manager only if your organization already has a large-scale deployment of Mesos, but it makes for a good cluster manager nonetheless
- Historically Mesos supported a variety of different modes (fine-grained and coarse-grained), but at this point, it supports only coarse-grained scheduling (fine-grained has been deprecated)
- Coarse-grained mode means that each Spark executor runs as a single Mesos task. Spark executors are sized according to the following application properties:
  - `spark.executor.memory`
  - `spark.executor.cores`
  - `spark.cores.max/spark.executor.cores`

## Submitting applications

- similar to doing so for Spark's other cluster managers
- For the most part you should favor cluster mode when using Mesos.
- Client mode requires some extra configuration on your part, especially with regard to distributing resources around the cluster
- in client mode, the driver needs extra configuration information in `spark-env.sh` to work with Mesos.

In `spark-env.sh` set some environment variables:

```
export MESOS_NATIVE_JAVA_LIBRARY=<path to libmesos.so>
```

This path is typically `<prefix>/lib/libmesos` so where the prefix is `/usr/local` by default

On Mac OS X, the library is called **libmesos.dylib** instead of **libmesos.so**:

```
export SPARK_EXECUTOR_URI=<URL of spark-2.2.0.tar.gz uploaded above>
```

Finally, set the Spark Application property **spark.executor.uri** to **<URL of spark2.2.0.tar.gz>**

Now, when starting a Spark application against the cluster, pass a **mesos://URL** as the master when creating a **SparkContext**, and set that property as a parameter in your **SparkConf** variable or the initialization of a **SparkSession**

```
// in Scala
import org.apache.spark.sql.SparkSession
val spark = SparkSession.builder
  .master("mesos://HOST:5050")
  .appName("my app")
  .config("spark.executor.uri", "<path to spark-2.2.0.tar.gz uploaded above>")
  .getOrCreate()
```

## Secure Deployment Configurations

- Spark also provides some low-level ability to make your applications run more securely, especially in untrusted environments
- majority of this setup will happen outside of Spark
- These configurations are primarily network-based to help Spark run in a more secure manner
- This means authentication, network encryption, and setting TLS and SSL configurations

## Cluster Networking Configurations

- Just as shuffles are important, there can be some things worth tuning on the network
- helpful when performing custom deployment configurations for your Spark clusters when you need to use proxies in between certain nodes

## Application Scheduling

- Spark has several facilities for scheduling resources between computations
- each Spark Application runs an independent set of executor processes
- Cluster managers provide the facilities for scheduling across Spark applications
- within each Spark application, multiple jobs (i.e., Spark actions) may be running concurrently if they were submitted by different threads
  - common if your application is serving requests over the network
  - Spark includes a fair scheduler to schedule resources within each application
- If multiple users need to share your cluster and run different Spark Applications, there are different options to manage allocation, depending on the cluster manager
- The simplest option, available on all cluster managers, is static partitioning of resources. With this approach, each application is given a maximum amount of resources that it can use, and holds onto those resources for the entire duration

- In **spark-submit** there are a number of properties that you can set to control the resource allocation of a particular application
- In addition, dynamic allocation (described next) can be turned on to let applications scale up and down dynamically based on their current number of pending tasks
- If, instead, you want users to be able to share memory and executor resources in a fine-grained manner, you can launch a single Spark Application and use thread scheduling within it to serve multiple requests in parallel

## Dynamic allocation

- If you would like to run multiple Spark Applications on the same cluster, Spark provides a mechanism to dynamically adjust the resources your application occupies based on the workload
- This means that your application can give resources back to the cluster if they are no longer used, and request them again later when there is demand
- useful if multiple applications share resources in your Spark cluster
- This feature is disabled by default and available on all coarse-grained cluster managers; that is, standalone mode, YARN mode, and Mesos coarse-grained mode
- There are two requirements for using this feature:
  - First, your application must set `spark.dynamicAllocation.enabled` to true
  - Second, you must set up an external shuffle service on each worker node in the same cluster and set `spark.shuffle.service.enabled` to true in your application
  - The purpose of the external shuffle service is to allow executors to be removed without deleting shuffle files written by them

## Miscellaneous Considerations

- number and type of applications you intend to be running
- YARN is great for HDFS-based applications but is not commonly used for much else
  - not well designed to support the cloud, because it expects information to be available on HDFS.
- Also, compute and storage is largely coupled together, meaning that scaling your cluster involves scaling both storage and compute instead of just one or the other
- Mesos does improve on this a bit conceptually, and it supports a wide range of application types, but it still requires pre-provisioning machines and, in some sense, requires buy-in at a much larger scale
- doesn't really make sense to have a Mesos cluster for only running Spark Applications
- Spark standalone mode is the lightest-weight cluster manager and is relatively simple to understand and take advantage of, but then you're going to be building more application management infrastructure that you could get much more easily by using YARN or Mesos
- Another challenge is managing different Spark versions
  - shouldn't try to run a variety of different applications running different Spark versions, and unless you use a well-managed service, you're going to need to spend a fair amount of time managing different setup scripts for different Spark services
- consider how you're going to set up logging, store logs for future reference, and allow end users to debug their applications



then applications

- These are more “out of the box” for YARN or Mesos and might need some tweaking if you’re using standalone
- One thing you might want to consider—or that might influence your decision making—is maintaining a **metastore** in order to maintain metadata about your stored datasets, such as a table catalog
- Maintaining an Apache Hive metastore, a topic beyond the scope of this book, might be something that’s worth doing to facilitate more productive, cross-application referencing to the same datasets.
- Depending on your workload, it might be worth considering using Spark’s external shuffle service.
- Typically Spark stores shuffle blocks (shuffle output) on a local disk on that particular node
- external shuffle service allows for storing those shuffle blocks so that they are available to all executors, meaning that you can arbitrarily kill executors and still have their shuffle outputs available to other applications
- need to configure at least some basic monitoring solution and help users debug their Spark jobs running on their clusters